
title: Build Work Estimate — Granite Peak Utilities -- outage triage agent author: Dewain Robinson
build_platform: github-copilot target_platform: copilot-studio estimate_id: est_20260903T000000_sample2
generated: 2026-09-03T00:00:00Z

Build Work Estimate — Granite Peak Utilities -- outage triage agent

Author: Dewain Robinson · **Estimate id:** est_20260903T000000_sample2 · **Generated:** 2026-09-03T00:00:00Z

SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE

This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.

It estimates the work of building only — never the cost of running what gets built. Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.

The figures are estimates and will be wrong. They are derived from historical consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.

Before any real budgeting use, this estimator must be modified for your organization — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.

Rates verified 2026-06-24 and change without notice. Calibration source: not applicable · Generated 2026-09-03T00:00:00Z

Estimate summary

What was estimated

INPUT	VALUE
Built with (development tool)	GitHub Copilot — metered in GitHub AI Credits
Built on (target environment)	Microsoft Copilot Studio — metered in Copilot Credits
Target harness	standard — does not bill for build or test
Licensing	Seat — GitHub Copilot Business, \$19.00/month x 3 seats = \$57.00/month
Evaluation cycles planned	3
Contingency reserve	30%
Calibration	Not applicable — the build platform is not metered from local session history

Totals

	BUILD — GITHUB AI CREDITS	TARGET — COPILOT CREDITS	COMBINED (USD)
Low	900	104	\$10.04
Likely	900	156	\$10.56
High	900	260	\$11.60
Likely + 30% reserve	1,170	203	\$13.73

Plan for \$10.56. Hold \$13.73 including the 30% reserve.

Two meters, not two options. The build and target figures are spent on the same project over the same period and add together; they are not alternatives to choose between.

- Copilot Credits are converted at \$0.01 each so the two meters can be shown in one column. They are separate budgets drawn on different accounts.
- The build figure is **notional** — on seat licensing no additional money is invoiced. See Licensing below for the share of the seat this build actually consumes.
- The build side has no measured spread of its own: the declared interaction volume is the estimate. Only the target side varies here.

Everything below explains how each of these figures was reached.

Platforms

	PLATFORM	METERED IN
Built with	GitHub Copilot	GitHub AI Credits
Built on	Microsoft Copilot Studio	Copilot Credits

These are different questions and different meters, and the same project spends on **both**. The building happens in a coding agent authoring the agent definition; the target platform is where it is deployed, previewed, evaluated and validated. Copilot Studio is a destination, not a build tool.


This is not a platform comparison tool. It reports one chosen build platform and one chosen target, each in its own meter. Platform decisions are not made on cost alone — capability, existing skills, governance, integration, and support all matter more than a build-time figure — and using this report to pick a platform would be using it for something it was not designed to answer.

Scope — this estimates the build, not the run

This figure covers the work of **building** the thing. It does not cover operating it afterwards.

OUT OF SCOPE	WHERE TO GO INSTEAD
Runtime / operational cost of what gets built	Copilot Studio agent usage estimator
End-user seat licences (M365 Copilot, Power Platform)	Your licensing/procurement process
Infrastructure, hosting, storage, egress	Your cloud cost tooling
Human labour — PM, design, QA time	Your delivery estimation process
Ongoing maintenance and support after delivery	A run cost, not a build cost

The build loop



This estimate plans **3 evaluation cycles**. Each cycle after the first adds **25%** of the original build back as remediation, giving a build-side multiplier of **1.50x**. An estimate that prices only the first pass is planning for a build where every evaluation passes first time.

Build cost — GitHub Copilot

Metered in **premium requests**. This is a different meter from Copilot Studio Copilot Credits, even though both are \$0.01 per credit.

Billing model: Legacy premium requests — Each interaction costs one premium request multiplied by the model's multiplier, drawn from a monthly plan allowance. Eligible Copilot Pro and Pro+ subscribers on existing annual plans remain on this model until their plan expires.

BUILD ACTIVITY	PREMIUM REQUESTS	BASIS
Premium requests	900	900 interactions x 1.00 model multiplier
Total	900	
Reserve (30%)	270	required contingency
Budget ask	1,170	

Against a monthly allowance of 1,500 premium requests, this build consumes **78%**.

Not metered

These consume no credits and are unlimited on paid plans, so they contribute nothing to this estimate however heavily they are used:

- Code completions
- Next edit suggestions

Rates verified 2026-09-03. Sources: [models and pricing](#) · [legacy premium requests](#) · [plans](#)

Target platform — Microsoft Copilot Studio

Target harness: `standard` — Standard harness bills after publish, and embedded test chat messages are not billed. Build-time credits are therefore near zero. Non-zero only where the build exercises billable side-effects (agent flow runs, AI Builder or content-processing calls) against a published agent.

On the standard harness, build and test in the interface are not billed. Billing starts after publish and embedded test chat does not count, so the work below consumes no credits. That is a correct result, not a missing one.

Had this been the GitHub Copilot harness, the same volume of preview and evaluation work would have cost roughly **6,048 credits (\$60.48)** — worth knowing before choosing a harness.

BUILD-AND-TEST ACTIVITY	CREDITS	AT \$0.01/CC	BASIS
Agent flow runs during build and test	156	\$1.56	1,200 actions across 3 cycles at 13 CC per 100
Total	156	\$1.56	
Reserve (30%)	47	\$0.47	required contingency
Budget ask	203	\$2.03	

The evaluation loop

Evaluation is not a one-off gate. Microsoft's guidance is an explicit cycle — define tests, run evaluations, analyse results, improve the agent, repeat — with a target pass rate of 80–90% and near 100% on core tests.

This estimate plans **3 cycles**: 315 evaluation runs in total, plus the remediation each failed cycle sends back to the build platform.

Velocity, not just cost. Evaluations are capped at **20 per agent node per day**, so this volume needs a minimum of **18 days** of elapsed time however much budget is available.

Human validation — a dependency, not a cost line

20.0 hours of interactive validation are planned in the Copilot Studio interface. Those hours are collected to size the test volume above — they are **not** estimated as labour cost, consistent with this tool metering agent consumption rather than people.

Someone must go into the interface, confirm behaviour, and make configuration changes between cycles. That step gates the loop, and no amount of build budget removes it.

Not included

- Production runtime once the agent is live
- Capacity pack sizing and overage enforcement
- End-user Microsoft 365 Copilot licence offsets

- Human hours for validation (collected to size test volume only)

For production runtime consumption use Microsoft's [agent usage estimator](#).

Rates verified 2026-09-03. Sources: [billing rates](#) · [harness billing](#) · [evaluation limits](#) · [iteration guidance](#)

Known limits

1. **Build only.** Nothing here is the cost of running the agent once it is live.
2. **Human validation is a dependency, not a line item.** Hours are collected to size test volume, never estimated as labour.
3. **Turn counts are the weakest input**, and thin buckets are flagged above.
4. **Rates go stale.** Re-verify against the sources cited.
5. **This is a sample.** Modify it for your organization before budgeting use.

Close the loop

```
python record_actual.py est_20260903T000000_sample2 --sessions <session-id> [...]
```

SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.