
title: Build Work Estimate — Granite Peak Utilities -- outage triage agent, built in Copilot Studio author: Dewain Robinson build_stack: copilot-studio estimate_id: est_20260903T093307_fc766b generated: 2026-09-03T09:33:07Z

Build Work Estimate — Granite Peak Utilities -- outage triage agent, built in Copilot Studio

Author: Dewain Robinson · **Estimate id:** est_20260903T093307_fc766b · **Generated:** 2026-09-03T09:33:07Z

SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE

This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.

It estimates the work of building only — never the cost of running what gets built. Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.

The figures are estimates and will be wrong. They are derived from historical consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.

Before any real budgeting use, this estimator must be modified for your organization — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.

Rates verified 2026-06-24 and change without notice. Calibration source: not applicable · Generated 2026-09-03T09:33:07Z

Build stack — Microsoft Copilot Studio

Metered in: Copilot Credits. Metered in Copilot Credits. Not tokens, and not GitHub AI Credits -- a separate meter that happens to share a rate.

The currency is decided by what you **build with**, not what you build **for**. Building a Microsoft workload with Claude Code is metered in tokens; building it in Copilot Studio is metered in Copilot Credits.

This is not a stack comparison tool. It reports one build stack, in that stack's own currency. Technology stack decisions are not made on cost alone — capability, existing skills, governance, integration, and support all matter more than a build-time figure — and using this report to pick a stack would be using it for something it was not designed to answer.

Scope — this estimates the build, not the run

This figure covers the work of **building** the thing. It does not cover operating it afterwards.

OUT OF SCOPE	WHERE TO GO INSTEAD
Runtime / operational cost of what gets built	Copilot Studio agent usage estimator
End-user seat licences (M365 Copilot, Power Platform)	Your licensing/procurement process
Infrastructure, hosting, storage, egress	Your cloud cost tooling
Human labour — PM, design, QA time	Your delivery estimation process
Ongoing maintenance and support after delivery	A run cost, not a build cost

Build-time Copilot Credits

These are the credits consumed **building** the agent. They say nothing about what it will cost once users start talking to it. For that, use Microsoft's [agent usage estimator](#).

Harness: `github-copilot` — GitHub Copilot harness bills from the moment you start building. Creating a solution with natural language, previewing, testing, and generating evaluations all consume credits. Credits cover LLM tokens, tools (knowledge and MCPs), and the harness itself.

BUILD ACTIVITY	CREDITS	AT \$0.01/CC	BASIS
Authoring / natural-language solution creation	6,720	\$67.20	160 turns x 4,200 tokens at premium tier
Preview and test iterations	540	\$5.40	45 runs x 6 interactions x 2 CC (Generative answer)
Preview and test model consumption	11,340	\$113.40	1,134,000 tokens at premium tier
Evaluation generation and runs	900	\$9.00	30 evals x 3,000 tokens at premium tier
Agent flow actions during build	117	\$1.17	900 actions at 13 CC per 100
Total build credits	19,617	\$196.17	
Reserve (30%)	5,885	\$58.85	required contingency
Budget ask	25,502	\$255.02	

Reasoning-model surcharge: 16,116 credits (\$161.16). Reasoning models bill the feature rate *plus* the premium token tier, so the effective tier is `premium` regardless of the `standard` tier selected.

Tier sensitivity

Tier selection is the highest-leverage variable in a Microsoft-side build estimate:

TIER	CC PER 1K TOKENS	\$ PER 1M TOKENS
basic	0.1	\$1.00
standard	1.5	\$15.00
premium ← in use	10.0	\$100.00

Not included

- Monthly production credit burn from end users
- Capacity pack sizing and overage enforcement
- Voice minutes
- End-user Microsoft 365 Copilot licence offsets
- Bring-your-own-model (including Azure Foundry), billed separately

Rates verified 2026-09-03. Sources: [billing_rates](#) · [pay-as-you-go](#) · [reasoning](#) · [harness](#)

Licensing

Consumption billing. Every unit consumed is billed, so the figures above are the expected charge.

Known limits

1. **Build only.** Says nothing about what the built thing costs to run.
2. **Rates go stale.** Verified 2026-09-03; re-verify against the sources above.
3. **This is a sample.** Modify it for your organization before budgeting use.

Close the loop

```
python record_actual.py est_20260903T093307_fc766b --sessions <session-id> [...]
```

SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.