

Build Work Estimate — Harbor Line Logistics -- dispatch modernization

Author: Dewain Robinson · **Estimate id:** est_20260903T091026_b91d9f · **Generated:** 2026-09-03T09:10:26Z

SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE

This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.

It estimates the work of building only — never the cost of running what gets built. Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.

The figures are estimates and will be wrong. They are derived from historical token consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.

Before any real budgeting use, this estimator must be modified for your organization — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.

Rates verified 2026-06-24 and change without notice. Calibration source: **measured** · Generated 2026-09-03T09:10:26Z

Scope — this estimates the build, not the run

This figure covers the agent turns consumed **building** the thing. It does not cover operating it afterwards.

OUT OF SCOPE	WHERE TO GO INSTEAD
Runtime / operational cost of what gets built	Copilot Studio agent usage estimator
End-user seat licences (M365 Copilot, Power Platform)	Your licensing/procurement process
Infrastructure, hosting, storage, egress	Your cloud cost tooling
Human labour — PM, design, QA time	Your delivery estimation process
Ongoing maintenance and support after delivery	A run cost, not a build cost

Summary

	AMOUNT
Base estimate	\$508.69
Reserve (25%)	\$127.17
Budget ask	\$635.86
Observed low	\$131.81
Observed high	\$2,989.13

Reserve adequacy

The reserve does not cover observed variance.

A **25%** reserve brings the budget ask to \$635.86. Comparable work in the calibration data has reached \$2,989.13. Full coverage would require **488%**.

This is the single most likely way this estimate causes a problem: not that the point figure is wrong, but that the contingency carried on top of it is too thin for how wide the real spread is.

Breakdown

ITEM	SIZE	FILES	UNKNOWN	TURNS	ESTIMATE	RANGE	N
Dispatch scheduling API	medium (brownfield)	11	2	696	\$256.20	\$56.45 – \$1,309.21	5
Carrier rate-card importer	small	4	1	165	\$60.72	\$33.12 – \$238.54	6
Legacy AS400 field mapping research	exploration (brownfield)	0	3	57	\$20.98	\$4.60 – \$859.50	7
Driver mobile handoff screens	medium	8	0	464	\$170.80	\$37.63 – \$581.87	5

Calibration basis

Measured from 24 local sessions spanning 2026-06-01 to 2026-08-30.

Cost per agent turn: **\$0.32**. Rates are **list price** — an organization on contracted rates must substitute its own.

Correction factors from recorded actuals

BUCKET	N	OBSERVED RATIO	APPLIED (SHRUNK)	IN EFFECT
medium	3	1.35x	1.18x	yes

Observed ratios are shrunk toward 1.0 in proportion to sample size, so a single recorded actual cannot swing future estimates.

Build-time Copilot Credits

These are the credits consumed **building** the agent. They say nothing about what it will cost once users start talking to it. For that, use Microsoft's [agent usage estimator](#).

Harness: `github-copilot` — GitHub Copilot harness bills from the moment you start building. Creating a solution with natural language, previewing, testing, and generating evaluations all consume credits. Credits cover LLM tokens, tools (knowledge and MCPs), and the harness itself.

BUILD ACTIVITY	CREDITS	AT \$0.01/CC	BASIS
Authoring / natural-language solution creation	6,300	\$63.00	140 turns x 4,500 tokens at premium tier
Preview and test iterations	420	\$4.20	35 runs x 6 interactions x 2 CC (Generative answer)
Preview and test token consumption	9,450	\$94.50	945,000 tokens at premium tier
Evaluation generation and runs	800	\$8.00	25 evals x 3,200 tokens at premium tier
Agent flow actions during build	78	\$0.78	600 actions at 13 CC per 100
Total build credits	17,048	\$170.48	
Reserve (25%)	4,262	\$42.62	required contingency
Budget ask	21,310	\$213.10	

Reasoning-model surcharge: 14,068 credits (\$140.68). Reasoning models bill the feature rate *plus* the premium token tier, so the effective tier is `premium` regardless of the `standard` tier selected.

Tier sensitivity

Tier selection is the highest-leverage variable in a Microsoft-side build estimate:

TIER	CC PER 1K TOKENS	\$ PER 1M TOKENS
basic	0.1	\$1.00
standard	1.5	\$15.00
premium ← in use	10.0	\$100.00

Not included

- Monthly production credit burn from end users
- Capacity pack sizing and overage enforcement
- Voice minutes
- End-user Microsoft 365 Copilot licence offsets
- Bring-your-own-model (including Azure Foundry), billed separately

Rates verified 2026-09-03. Sources: [billing_rates](#) · [pay-as-you-go](#) · [reasoning](#) · [harness](#)

Assumptions and sensitivities

- Turn counts come from the calibration profile, and are the weakest input in this model. Where `n` is small, treat the figure as indicative only.
- Cost scales with context size and turn count. Session hygiene — clearing between unrelated tasks — moves the number more than most scope changes do.
- Brownfield work is scaled by 1.5x over greenfield.
- Each declared unknown widens the upper bound by 25%.

Known limits

1. **Build only.** Says nothing about what the built thing costs to run.
2. **Turn counts are the weakest input**, and thin buckets are flagged above.
3. **List pricing only.** Contracted rates must be substituted.
4. **This machine only.** Sessions from other devices are not visible.
5. **Rates go stale.** Both providers change pricing without notice.

Close the loop

When this build is done, record what it actually cost. That is how the estimator gets better — and without it, it cannot.

```
python record_actual.py est_20260903T091026_b91d9f --sessions <session-id> [...]
```

SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.