

title: Build Work Estimate — Harbor Line Logistics -- dispatch agent author: Dewain Robinson
build_platform: claude-code target_platform: copilot-studio estimate_id: est_20260903T100519_a78168
generated: 2026-09-03T10:05:19Z

Build Work Estimate — Harbor Line Logistics -- dispatch agent

Author: Dewain Robinson · Estimate id: est_20260903T100519_a78168 · Generated: 2026-09-03T10:05:19Z

SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE

This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.

It estimates the work of building only — never the cost of running what gets built. Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.

The figures are estimates and will be wrong. They are derived from historical consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.

Before any real budgeting use, this estimator must be modified for your organization — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.

Rates verified 2026-06-24 and change without notice. Calibration source: measured · Generated 2026-09-03T10:05:19Z

Platforms

	PLATFORM	METERED IN
Built with	Claude Code	USD (tokens)
Built on	Microsoft Copilot Studio	Copilot Credits

These are different questions and different meters, and the same project spends on **both**. The building happens in a coding agent authoring the agent definition; the target platform is where it is deployed, previewed, evaluated and validated. Copilot Studio is a destination, not a build tool.

This is not a platform comparison tool. It reports one chosen build platform and one chosen target, each in its own meter. Platform decisions are not made on cost alone — capability, existing skills, governance, integration, and support all matter more than a build-time figure — and using this report to pick a platform would be using it for something it was not designed to answer.

Scope — this estimates the build, not the run

This figure covers the work of **building** the thing. It does not cover operating it afterwards.


OUT OF SCOPE	WHERE TO GO INSTEAD
Runtime / operational cost of what gets built	Copilot Studio agent usage estimator
End-user seat licences (M365 Copilot, Power Platform)	Your licensing/procurement process
Infrastructure, hosting, storage, egress	Your cloud cost tooling
Human labour — PM, design, QA time	Your delivery estimation process
Ongoing maintenance and support after delivery	A run cost, not a build cost

Summary

COMPONENT	METER	TOTAL	WITH 25% RESERVE
Build — authoring and remediation	USD (tokens)	\$890.21	\$1,112.76
Target — preview, test and evaluation	Copilot Credits	32,290	40,362

Two meters, not two options. The figures above are spent on the same project over the same period and add together; they are not alternatives to choose between.

The build loop



```
remediate          <-- evaluations fail
(repeat)
```

This estimate plans **4 evaluation cycles**. Each cycle after the first adds **25%** of the original build back as remediation, giving a build-side multiplier of **1.75x**. An estimate that prices only the first pass is planning for a build where every evaluation passes first time.

Build — Claude Code

	AMOUNT
Base estimate (incl. remediation)	\$890.21
Reserve (25%)	\$222.55
Budget ask	\$1,112.76
Observed low	\$230.67
Observed high	\$5,230.98

Reserve adequacy

The reserve does not cover observed variance.

A **25%** reserve reaches \$1,112.76. Comparable work has reached \$5,230.98. Full coverage would require **488%**.

ITEM	SIZE	FILES	TURNS	ESTIMATE	RANGE	N
Agent instructions, topics and tool definitions	medium (brownfield)	11	1,218	\$448.35	\$98.79 – \$2,291.12	5
Dataverse schema and rate-card importer	small	4	289	\$106.26	\$57.96 – \$417.44	6
Legacy AS400 field mapping research	exploration (brownfield)	0	100	\$36.72	\$8.05 – \$1,504.12	7
Driver mobile handoff screens	medium	8	812	\$298.90	\$65.85 – \$1,018.27	5

Calibration basis

Measured from 24 local sessions spanning 2026-06-01 to 2026-08-30.

Cost per agent turn: **\$0.32**, at **list price** — an organization on contracted rates must substitute its own.

Licensing

Seat-based licensing (Claude Max). Usage draws on an allowance already paid for, so no additional money changes hands for this build -- but the seat is not free.

Share of a typical month's allowance	188%
Seat cost per month	\$200.00
Attributable cost of this build	\$376.15
Already committed to other work	45%
Total committed	233%

The attributable cost is the seat's monthly price apportioned by the share of the allowance this build consumes. Nothing extra is invoiced, but this is the real cost of the capacity the build uses up.

Allowance overrun

This build plus existing workload exceeds the allowance by 133%.

Work will stall at the limit unless overage is enabled, at which point it bills on top of the seat. Overage exposure at the same rate: **\$266.15**.

Options: spread the build across allowance periods, move part of it to consumption billing, add seats, or reduce the other committed work.

Window risk

Seat allowances refill on **5-hour rolling and weekly** windows, not only monthly. A build that fits comfortably in a month can still exhaust a short window and stall.

This build is marked as concentrated — compressed into a short period. Monthly headroom will not protect it; expect to hit shorter windows and plan for pauses.

Notional value at list rates: **\$890.21** — what this build would cost on consumption billing. Shown for scale; it is not a charge.

Target platform — Microsoft Copilot Studio

Target harness: `github-copilot` — GitHub Copilot harness bills from the moment you start building. Creating a solution with natural language, previewing, testing, and generating evaluations all consume credits. Credits cover LLM tokens, tools (knowledge and MCPs), and the harness itself.

BUILD-AND-TEST ACTIVITY	CREDITS	AT \$0.01/CC	BASIS
Interactive validation in the Copilot Studio interface	16,800	\$168.00	16.0 human hours x 25 interactions/hour = 400 interactions
Evaluation runs	15,360	\$153.60	40 cases x 3 repeats x 4 cycles = 480 runs
Agent flow runs during build and test	130	\$1.30	1,000 actions across 4 cycles at 13 CC per 100
Total	32,290	\$322.90	
Reserve (25%)	8,072	\$80.73	required contingency
Budget ask	40,362	\$403.62	

Reasoning-model surcharge: 25,840 credits (\$258.40). Reasoning models bill the feature rate *plus* the premium token tier, so the effective tier is `premium` whatever the `standard` tier selected.

The evaluation loop

Evaluation is not a one-off gate. Microsoft's guidance is an explicit cycle — define tests, run evaluations, analyse results, improve the agent, repeat — with a target pass rate of 80–90% and near 100% on core tests.

This estimate plans **4 cycles**: 480 evaluation runs in total, plus the remediation each failed cycle sends back to the build platform.

Velocity, not just cost. Evaluations are capped at **20 per agent node per day**, so this volume needs a minimum of **24 days** of elapsed time however much budget is available.

Human validation — a dependency, not a cost line

16.0 hours of interactive validation are planned in the Copilot Studio interface. Those hours are collected to size the test volume above — they are **not** estimated as labour cost, consistent with this tool metering agent consumption rather than people.

Someone must go into the interface, confirm behaviour, and make configuration changes between cycles. That step gates the loop, and no amount of build budget removes it.

Not included

- Production runtime once the agent is live
- Capacity pack sizing and overage enforcement
- End-user Microsoft 365 Copilot licence offsets
- Human hours for validation (collected to size test volume only)

For production runtime consumption use Microsoft's [agent usage estimator](#).

Rates verified 2026-09-03. Sources: [billing rates](#) · [harness billing](#) · [evaluation limits](#) · [iteration guidance](#)

Known limits

1. **Build only.** Nothing here is the cost of running the agent once it is live.
2. **Human validation is a dependency, not a line item.** Hours are collected to size test volume, never estimated as labour.
3. **Turn counts are the weakest input**, and thin buckets are flagged above.
4. **Rates go stale.** Re-verify against the sources cited.
5. **This is a sample.** Modify it for your organization before budgeting use.

Close the loop

```
python record_actual.py est_20260903T100519_a78168 --sessions <session-id> [...]
```

SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.