---

title: Build Work Estimate — Harbor Line Logistics -- Copilot Studio dispatch agent, built in Claude Code author: Dewain Robinson build_stack: claude-code estimate_id: est_20260903T093307_718b2e generated: 2026-09-03T09:33:07Z

Build Work Estimate — Harbor Line Logistics -- Copilot Studio dispatch agent, built in Claude Code

Author: Dewain Robinson · **Estimate id:** `est_20260903T093307_718b2e` · **Generated:** 2026-09-03T09:33:07Z

> **SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE**
>
> This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.
>
> **It estimates the work of building only — never the cost of running what gets built.** Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.
>
> **The figures are estimates and will be wrong.** They are derived from historical consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.
>
> **Before any real budgeting use, this estimator must be modified for your organization** — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.
>
> Rates verified 2026-06-24 and change without notice. Calibration source: `measured` · Generated 2026-09-03T09:33:07Z

Build stack — Claude Code

Metered in: USD (tokens). Metered in input/output/cache tokens, priced per model.

> The currency is decided by what you **build with**, not what you build **for**. Building a Microsoft workload with Claude Code is metered in tokens; building it in Copilot Studio is metered in Copilot Credits.
>
> **This is not a stack comparison tool.** It reports one build stack, in that stack's own currency. Technology stack decisions are not made on cost alone — capability, existing skills, governance, integration, and support all matter more than a build-time figure — and using this report to pick a stack would be using it for something it was not designed to answer.

Scope — this estimates the build, not the run

This figure covers the work of **building** the thing. It does not cover operating it afterwards.

OUT OF SCOPE	WHERE TO GO INSTEAD
Runtime / operational cost of what gets built	Copilot Studio agent usage estimator
End-user seat licences (M365 Copilot, Power Platform)	Your licensing/procurement process
Infrastructure, hosting, storage, egress	Your cloud cost tooling
Human labour — PM, design, QA time	Your delivery estimation process
Ongoing maintenance and support after delivery	A run cost, not a build cost

Summary

	AMOUNT
Base estimate	$508.69
Reserve (25%)	$127.17
Budget ask	**$635.86**
Observed low	$131.81
Observed high	$2,989.13

Reserve adequacy

The reserve does not cover observed variance.

A **25%** reserve brings the budget ask to $635.86. Comparable work in the calibration data has reached $2,989.13. Full coverage would require **488%**.

This is the single most likely way this estimate causes a problem: not that the point figure is wrong, but that the contingency carried on top of it is too thin for how wide the real spread is.

Breakdown

ITEM	SIZE	FILES	UNKNOWNS	TURNS	ESTIMATE	RANGE	N
Agent topic + tool authoring	medium (brownfield)	11	2	696	$256.20	$56.45 – $1,309.21	5
Dataverse schema and rate-card importer	small	4	1	165	$60.72	$33.12 – $238.54	6
Legacy AS400 field mapping research	exploration (brownfield)	0	3	57	$20.98	$4.60 – $859.50	7
Driver mobile handoff screens	medium	8	0	464	$170.80	$37.63 – $581.87	5

Calibration basis

Measured from 24 local sessions spanning 2026-06-01 to 2026-08-30.

Cost per agent turn: **$0.32**. Rates are **list price** — an organization on contracted rates must substitute its own.

Correction factors from recorded actuals

BUCKET	N	OBSERVED RATIO	APPLIED (SHRUNK)	IN EFFECT
medium	3	1.35x	1.18x	yes

Observed ratios are shrunk toward 1.0 in proportion to sample size, so a single recorded actual cannot swing future estimates.

Licensing

Seat-based licensing (Claude Max). Usage draws on an allowance already paid for, so no additional money changes hands for this build -- but the seat is not free.

Share of a typical month's allowance	**107%**
Seat cost per month	$200.00
Attributable cost of this build	**$214.94**
Already committed to other work	45%
Total committed	**152%**

The attributable cost is the seat's monthly price apportioned by the share of the allowance this build consumes. Nothing extra is invoiced, but this is the real cost of the capacity the build uses up.

Allowance overrun

This build plus existing workload exceeds the allowance by 52%.

Work will stall at the limit unless overage is enabled, at which point it bills on top of the seat. Overage exposure at the same rate: **$104.94**.

Options: spread the build across allowance periods, move part of it to consumption billing, add seats, or reduce the other committed work.

Window risk

Seat allowances refill on **5-hour rolling and weekly** windows, not only monthly. A build that fits comfortably in a month can still exhaust a short window and stall.

This build is marked as concentrated — compressed into a short period. Monthly headroom will not protect it; expect to hit shorter windows and plan for pauses.

Notional value at list rates: **$508.69** — what this build would cost on consumption billing. Shown for scale; it is not a charge.

Assumptions and sensitivities

- Turn counts come from the calibration profile, and are the weakest input in this model. Where `n` is small, treat the figure as indicative only.
- Cost scales with context size and turn count. Session hygiene — clearing between unrelated tasks — moves the number more than most scope changes do.
- Brownfield work is scaled by 1.5x over greenfield.
- Each declared unknown widens the upper bound by 25%.

Known limits

1. **Build only.** Says nothing about what the built thing costs to run.
2. **Turn counts are the weakest input**, and thin buckets are flagged above.
3. **List pricing only.** Contracted rates must be substituted.
4. **This machine only.** Sessions from other devices are not visible.
5. **Rates go stale.** Both providers change pricing without notice.

Close the loop

When this build is done, record what it actually cost. That is how the estimator gets better — and without it, it cannot.

```
python record_actual.py est_20260903T093307_718b2e --sessions <session-id> [...]
```

SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.