

---

title: Build Work Estimate — Kestrel Financial Group -- Advisor Intelligence Platform (phases 1-2) author: Dewain Robinson build\_platform: github-copilot target\_platform: both estimate\_id: est\_20260903T000000\_kestrelgh generated: 2026-09-03T00:00:00Z

---

# Build Work Estimate — Kestrel Financial Group

## -- Advisor Intelligence Platform (phases 1-2)

**Author:** Dewain Robinson · **Estimate id:** est\_20260903T000000\_kestrelgh · **Generated:** 2026-09-03T00:00:00Z

### SAMPLE — BUILD ESTIMATE ONLY, NOT A QUOTE

This document is generated by a **sample estimator** published as a demonstration of *how* an organization could build one. It is not a quote, a bid, a budget authority, or a commitment of any kind.

**It estimates the work of building only — never the cost of running what gets built.** Runtime, end-user licences, infrastructure, and human labour are all outside its scope and are not represented in any figure below.

**The figures are estimates and will be wrong.** They are derived from historical consumption patterns that may not resemble the work being estimated. Observed cost for comparable work in the calibration data spans more than 100x between the cheapest and most expensive instances. Treat the range as the estimate; treat the point figure as the midpoint of a guess.

**Before any real budgeting use, this estimator must be modified for your organization** — recalibrated against your own usage history, repriced against your contracted rates rather than list price, and adjusted for your own delivery patterns, model choices, and review overhead.

Rates verified 2026-09-03 and change without notice. Calibration source: measured · Generated 2026-09-03T00:00:00Z

# Estimate summary

## What was estimated

| INPUT                         | VALUE                                                                                                   |
|-------------------------------|---------------------------------------------------------------------------------------------------------|
| Built with (development tool) | GitHub Copilot — metered in GitHub AI Credits                                                           |
| Built on (target environment) | Copilot Studio and Azure — metered in Copilot Credits + Azure consumption                               |
| Target harness                | github-copilot — bills for build, test and evaluation                                                   |
| Licensing                     | Seat — GitHub Copilot Enterprise, \$39.00/month x 6 seats x 5 months = <b>\$1,170.00</b> over the build |
| Evaluation cycles planned     | 6                                                                                                       |
| Contingency reserve           | 30%                                                                                                     |
| Work sizing                   | Turn counts from measured Claude Code calibration (24 sessions), used as a work-size proxy              |

## Totals

|                      | BUILD — GITHUB AI CREDITS | TARGET — COPILOT CREDITS | COMBINED (USD)     |
|----------------------|---------------------------|--------------------------|--------------------|
| Low                  | 33,960                    | 416,820                  | \$20,607.80        |
| Likely               | 101,853                   | 480,024                  | <b>\$21,918.77</b> |
| High                 | 581,476                   | 606,432                  | \$27,979.08        |
| Likely + 30% reserve | 132,408                   | 624,031                  | <b>\$28,494.39</b> |

Plan for \$21,918.77. Hold \$28,494.39 including the 30% reserve.

**Two meters, not two options.** The build and target figures are spent on the same project over the same period and add together; they are not alternatives to choose between.

- Copilot Credits are converted at \$0.01 each so the two meters can be shown in one column. They are separate budgets drawn on different accounts.
- The build figure is **notional** — on seat licensing no additional money is invoiced. See Licensing below for the share of the seat this build actually consumes.

- The build side is sized from the same work breakdown using turn medians measured on Claude Code. Turn count is treated as a property of the work rather than of the tool; the price per turn is not. That is an assumption, not a measurement of GitHub Copilot.

Everything below explains how each of these figures was reached.

## Platforms

|            | PLATFORM                 | METERED IN                          |
|------------|--------------------------|-------------------------------------|
| Built with | GitHub Copilot           | GitHub AI Credits                   |
| Built on   | Copilot Studio and Azure | Copilot Credits + Azure consumption |

These are different questions and different meters, and the same project spends on **both**. The building happens in a coding agent authoring the agent definition; the target platform is where it is deployed, previewed, evaluated and validated. Copilot Studio is a destination, not a build tool.

**This is not a platform comparison tool.** It reports one chosen build platform and one chosen target, each in its own meter. Platform decisions are not made on cost alone — capability, existing skills, governance, integration, and support all matter more than a build-time figure — and using this report to pick a platform would be using it for something it was not designed to answer.

## Scope — this estimates the build, not the run

This figure covers the work of **building** the thing. It does not cover operating it afterwards.

| OUT OF SCOPE                                          | WHERE TO GO INSTEAD                                  |
|-------------------------------------------------------|------------------------------------------------------|
| Runtime / operational cost of what gets built         | <a href="#">Copilot Studio agent usage estimator</a> |
| End-user seat licences (M365 Copilot, Power Platform) | Your licensing/procurement process                   |
| Infrastructure, hosting, storage, egress              | Your cloud cost tooling                              |
| Human labour — PM, design, QA time                    | Your delivery estimation process                     |
| Ongoing maintenance and support after delivery        | A run cost, not a build cost                         |

## The build loop

|                |                          |
|----------------|--------------------------|
| GitHub Copilot | Copilot Studio and Azure |
| -----          | -----                    |

```

author definition --> deploy
                    preview / interactive test <- human
                    run evaluations
remediate          <-- evaluations fail
(repeat)

```

This estimate plans **6 evaluation cycles**. Each cycle after the first adds **25%** of the original build back as remediation, giving a build-side multiplier of **2.25x**. An estimate that prices only the first pass is planning for a build where every evaluation passes first time.

## Build cost — GitHub Copilot

Metered in **GitHub AI Credits**. This is a different meter from Copilot Studio Copilot Credits, even though both are \$0.01 per credit.

**Billing model:** Usage-based GitHub AI Credits — Interactions consume input, output, and cached tokens. GitHub prices those at the model's published rates and converts the result to AI Credits at \$0.01 per credit. On Business and Enterprise plans credits are pooled at the billing-entity level.

| BUILD ACTIVITY              | GITHUB AI CREDITS | BASIS                                                                                                 |
|-----------------------------|-------------------|-------------------------------------------------------------------------------------------------------|
| Chat and agent interactions | 101,853           | 23,577 interactions, 141,462,000 tokens (15% output) at \$5.00/\$25.00 per 1M, less 10% Auto discount |
| <b>Total</b>                | <b>101,853</b>    |                                                                                                       |
| Reserve (30%)               | 30,556            | required contingency                                                                                  |
| <b>Budget ask</b>           | <b>132,408</b>    |                                                                                                       |

At \$0.01 per credit that is **\$1,018.53** total, **\$1,324.08** with reserve.

### Not metered

These consume no credits and are unlimited on paid plans, so they contribute nothing to this estimate however heavily they are used:

- Code completions
- Next edit suggestions

Rates verified 2026-09-03. Sources: [models and pricing](#) · [legacy premium requests](#) · [plans](#)

# Target platform — Microsoft Copilot Studio

**Target harness:** `github-copilot` — GitHub Copilot harness bills from the moment you start building. Creating a solution with natural language, previewing, testing, and generating evaluations all consume credits. Credits cover LLM tokens, tools (knowledge and MCPs), and the harness itself.

| BUILD-AND-TEST ACTIVITY                                | CREDITS        | AT<br>\$0.01/CC   | BASIS                                                                                                |
|--------------------------------------------------------|----------------|-------------------|------------------------------------------------------------------------------------------------------|
| Interactive validation in the Copilot Studio interface | 100,800        | \$1,008.00        | 96.0 human hours x 25 interactions/hour = 2,400 interactions                                         |
| Evaluation runs                                        | 378,288        | \$3,782.88        | 568 cases x 3 repeats x 6 cycles = 10,224 runs                                                       |
| Agent flow runs during build and test                  | 936            | \$9.36            | 7,200 actions across 6 cycles at 13 CC per 100 -- test runs in the designer and test chat are exempt |
| <b>Total</b>                                           | <b>480,024</b> | <b>\$4,800.24</b> |                                                                                                      |
| Reserve (30%)                                          | 144,007        | \$1,440.07        | required contingency                                                                                 |
| <b>Budget ask</b>                                      | <b>624,031</b> | <b>\$6,240.31</b> |                                                                                                      |

**Reasoning-model surcharge: 385,764 credits (\$3,857.64).** Reasoning models bill the feature rate *plus* the premium token tier, so the effective tier is `premium` whatever the `standard` tier selected.

## The evaluation loop

Evaluation is not a one-off gate. Microsoft's guidance is an explicit cycle — define tests, run evaluations, analyse results, improve the agent, repeat — with a target pass rate of 80–90% and near 100% on core tests.

This estimate plans **6 cycles**: 10,224 evaluation runs in total, plus the remediation each failed cycle sends back to the build platform.

**Velocity, not just cost.** Evaluations are capped at **20 per agent node per day**, so this volume needs a minimum of **516 days** of elapsed time however much budget is available.

## Human validation — a dependency, not a cost line

**96.0 hours** of interactive validation are planned in the Copilot Studio interface. Those hours are collected to size the test volume above — they are **not** estimated as labour cost, consistent with this tool metering agent consumption rather than people.

Someone must go into the interface, confirm behaviour, and make configuration changes between cycles. That step gates the loop, and no amount of build budget removes it.

Not included

- Production runtime once the agent is live
- Agent flow test runs in the designer or test chat (documented exemption)
- Capacity pack sizing and overage enforcement
- End-user Microsoft 365 Copilot licence offsets
- Human hours for validation (collected to size test volume only)

For production runtime consumption use Microsoft's [agent usage estimator](#).

Rates verified 2026-09-03. Sources: [billing rates](#) · [harness billing](#) · [evaluation limits](#) · [iteration guidance](#)

Target platform — Azure

Azure consumption during build and test, as supplied. Azure rates are not bundled here — they depend entirely on the services chosen.

| ITEM                                         | USD         |
|----------------------------------------------|-------------|
| App Service and Functions (dev/test/staging) | \$2,400.00  |
| Cosmos DB (three environments)               | \$3,100.00  |
| Azure AI Search with semantic ranking        | \$4,200.00  |
| API Management                               | \$2,800.00  |
| Service Bus, Storage, Key Vault              | \$900.00    |
| Observability ingestion                      | \$1,600.00  |
| Front Door and networking                    | \$1,100.00  |
| Reserve (30%)                                | \$4,830.00  |
| Budget ask                                   | \$20,930.00 |

Known limits

1. **Build only.** Nothing here is the cost of running the agent once it is live.
2. **Human validation is a dependency, not a line item.** Hours are collected to size test volume, never estimated as labour.

3. **Turn counts are the weakest input**, and thin buckets are flagged above.
4. **Rates go stale.** Re-verify against the sources cited.
5. **This is a sample.** Modify it for your organization before budgeting use.

## Close the loop

---

```
python record_actual.py est_20260903T000000_kestrelgh --sessions <session-id> [...]
```

---

*SAMPLE — build estimate only, not a quote. Excludes runtime cost. Modify for your organization before budgeting use.*